



Chat with an OpenAI model



Source: [R/provider-openai.R](#)

This is the main interface to [OpenAI](#)'s models, using the **responses API**. You can use this to access OpenAI's latest models and features like image generation and web search. If you need to use an OpenAI-compatible API from another provider, or the **chat completions API** with OpenAI, use [`chat_openai_compatible\(\)`](#) instead.

Note that a ChatGPT Plus membership does not grant access to the API. You will need to sign up for a developer account (and pay for it) at the [developer platform](#).

Usage

```
chat_openai(  
  system_prompt = NULL,  
  base_url = "https://api.openai.com/v1",  
  api_key = NULL,  
  credentials = NULL,  
  model = NULL,  
  params = NULL,  
  api_args = list(),  
  api_headers = character(),  
  service_tier = c("auto", "default", "flex", "prior",  
  echo = c("none", "output", "all"))  
)  
  
models_openai(
```

)

Arguments

system_prompt

A system prompt to set the behavior of the assistant.

base_url

The base URL to the endpoint; the default is OpenAI's public API.

api_key

lifecycle deprecated Use `credentials` instead.



Override the default credentials. You generally should not need this argument; instead set the `OPENAI_API_KEY` environment variable. The best place to set this is in

`.Renvironment`, which you can easily edit by calling `usethis::edit_r_environ()`.

If you do need additional control, this argument takes a zero-argument function that returns either a string (the API key), or a named list (added as additional headers to every request).

model

The model to use for the chat (defaults to "gpt-4.1"). We regularly update the default, so we strongly recommend explicitly specifying a model for anything other than casual use. Use `models_openai()` to see all options.

params

api_args

Named list of arbitrary extra arguments appended to the body of every chat API call. Combined with the body object generated by ellmer with [modifyList\(\)](#).

api_headers

Named character vector of arbitrary extra headers appended to every chat API call.

service_tier

Request a specific service tier. There are four options:

- "auto" (default): uses the service tier configured in Project settings.
- "default": standard pricing and performance.
- "flex": slower and cheaper.
- "priority": faster and more expensive.

echo

One of the following options:

- none : don't emit any output (default when running in a function).
- output : echo text and tool-calling output as it streams in (default when running at the console).
- all : echo all input and output.

Note this only affects the [chat\(\)](#) method.

Value

A [Chat](#) object.

```
chat_aws_bedrock(), chat_azure_openai(),
chat_cloudflare(), chat_databricks(),
chat_deepseek(), chat_github(),
chat_google_gemini(), chat_groq(),
chat_huggingface(), chat_mistral(),
chat_ollama(), chat_openai_compatible(),
chat_openrouter(), chat_perplexity(),
chat_portkey()
```

Examples

```
chat <- chat_openai()
#> Using model = "gpt-4.1".
chat$chat("
  What is the difference between a tibble and a data frame?
  Answer with a bulleted list
")
#> - **Origin:**
#>   - **Data Frame:** Base R object for storing tabular data.
#>   - **Tibble:** Modern reimagining of data frames from the tidyverse (tibble package).
#>
#> - **Printing:**
#>   - **Data Frame:** Prints the entire object, possibly to the console.
#>   - **Tibble:** Prints a preview (first 10 rows by default) to the screen, making output more readable.
#>
#> - **Subsetting:**
#>   - **Data Frame:** May simplify to a vector when subsetting by a single column.
#>   - **Tibble:** Always returns a tibble when subsetting by a single column.
#>   - **Column Names:** 
#>     - **Data Frame:** Allows non-syntactic names (variable names starting with numbers or underscores).
```

ellmer 0.4.0

```
#>
#> - **Data Types:**  
#>   - **Data Frame:** Converts strings to factors by default  
#>   (unless otherwise specified).  
#>   - **Tibble:** Does not convert strings to factors.  
#>  
#> - **Partial Matching:**  
#>   - **Data Frame:** Allows partial matching of column names.  
#>   - **Tibble:** Does not allow partial matching of column  
#>   names.  
#>  
#> - **Performance:**  
#>   - **Data Frame:** Slightly faster for basic operations on  
#>   simpler structures.  
#>   - **Tibble:** Slightly slower but offers better performance  
#>   especially for big data analysis.  
#>  
#> - **Use in Tidyverse:**  
#>   - **Data Frame:** Used in base R workflows.  
#>   - **Tibble:** Default in tidyverse packages (ggplot2, dplyr, etc.).  
#>  
#> - **Row Names:**  
#>   - **Data Frame:** Supports row names.  
#>   - **Tibble:** Does not support row names; store them if needed.  
#>  
#> In summary, tibbles are a modern and tidyverse-friendly replacement for data frames with improved usability and printing.
```



```
chat$chat("Tell me three funny jokes about statisticians")  
#> Absolutely! Here are three funny jokes about statisticians:  
#>  
#> 1. **Why did the statistician bring a ladder to a party?**  
#>    Because they heard the drinks were on the house!  
#>  
#> 2. **How does a statistician catch a lion?**  
#>    They build a cage, label it as “population,” and then wait  
#>    for the lion to walk right in—it’s all about sampling!  
#>
```

ON THIS PAGE

[Usage](#)

[Arguments](#)

[Value](#)

[See also](#)

[Examples](#)